

```

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import warnings
warnings.filterwarnings('ignore')
import seaborn as sns

Customers = pd.read_excel('C:\\Users\\hp\\OneDrive\\Desktop\\UNIFIED
PROJECTS\\Budget Sales data\\AdventureWorks_Database.xlsx',
                        'Customers',
                        dtype={'CustomerKey':str},

parse_dates=['BirthDate', 'DateFirstPurchase']
)

Product = pd.read_excel('C:\\Users\\hp\\OneDrive\\Desktop\\UNIFIED
PROJECTS\\Budget Sales data\\AdventureWorks_Database.xlsx',
                        'Product',
                        dtype={'ProductKey':str},
                        parse_dates=['StartDate']
)

Sales = pd.read_excel('C:\\Users\\hp\\OneDrive\\Desktop\\UNIFIED
PROJECTS\\Budget Sales data\\AdventureWorks_Database.xlsx',
                        'Sales',
                        dtype={'ProductKey':str,
                                'CustomerKey':str,
                                'PromotionKey':str,
                                'SalesTerritoryKey':str},
                        parse_dates=['OrderDate', 'ShipDate']
)

Sales['DateKey'] = Sales['OrderDate'].astype(str)

Territory = pd.read_excel('C:\\Users\\hp\\OneDrive\\Desktop\\UNIFIED
PROJECTS\\Budget Sales data\\AdventureWorks_Database.xlsx',
                        'Territory',
                        dtype={'SalesTerritoryKey':str}
)

temp_data = pd.merge(Sales, Product, on='ProductKey', how='inner')
df = pd.merge(temp_data, Customers, on='CustomerKey', how='inner')
df = pd.merge(df, Territory, on='SalesTerritoryKey', how='inner')

df.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 58189 entries, 0 to 58188
Data columns (total 58 columns):
#   Column                Non-Null Count  Dtype
---  -
0   ProductKey            58189 non-null  object

```

1	OrderDate	58189	non-null	datetime64[ns]
2	ShipDate	58189	non-null	datetime64[ns]
3	CustomerKey	58189	non-null	object
4	PromotionKey	58189	non-null	object
5	SalesTerritoryKey	58189	non-null	object
6	SalesOrderNumber	58189	non-null	object
7	SalesOrderLineNumber	58189	non-null	int64
8	OrderQuantity	58189	non-null	int64
9	UnitPrice	58189	non-null	float64
10	TotalProductCost	58189	non-null	float64
11	SalesAmount	58189	non-null	float64
12	TaxAmt	58189	non-null	float64
13	Unnamed: 13	0	non-null	float64
14	Unnamed: 14	0	non-null	float64
15	Unnamed: 15	58189	non-null	float64
16	Unnamed: 16	58189	non-null	float64
17	Unnamed: 17	0	non-null	float64
18	Unnamed: 18	58189	non-null	float64
19	Unnamed: 19	0	non-null	float64
20	StandardCost_x	58189	non-null	float64
21	List Price	58189	non-null	float64
22	Unnamed: 22	0	non-null	float64
23	diif std cost	58189	non-null	int64
24	diff list price	58189	non-null	int64
25	DateKey	58189	non-null	object
26	ProductName	58189	non-null	object
27	SubCategory	58189	non-null	object
28	Category	58189	non-null	object
29	StandardCost_y	58189	non-null	float64
30	Color	30747	non-null	object
31	ListPrice	58189	non-null	float64
32	DaysToManufacture	58189	non-null	int64
33	ProductLine	58189	non-null	object
34	ModelName	58189	non-null	object
35	Photo	58189	non-null	object
36	ProductDescription	58189	non-null	object
37	StartDate	58189	non-null	datetime64[ns]
38	FirstName	58189	non-null	object
39	LastName	58189	non-null	object
40	FullName	58189	non-null	object
41	BirthDate	58189	non-null	datetime64[ns]
42	MaritalStatus	58189	non-null	object
43	Gender	58189	non-null	object
44	YearlyIncome	58189	non-null	int64
45	TotalChildren	58189	non-null	int64
46	NumberChildrenAtHome	58189	non-null	int64
47	Education	58189	non-null	object
48	Occupation	58189	non-null	object
49	HouseOwnerFlag	58189	non-null	int64

```

50  NumberCarsOwned      58189 non-null int64
51  AddressLine1         58189 non-null object
52  DateFirstPurchase    58189 non-null datetime64[ns]
53  CommuteDistance      58189 non-null object
54  Region               58189 non-null object
55  Country              58189 non-null object
56  Group                58189 non-null object
57  RegionImage          58189 non-null object
dtypes: datetime64[ns](5), float64(16), int64(10), object(27)
memory usage: 26.2+ MB

```

```
df.describe().transpose()
```

	count	mean	std	min
\				
SalesOrderLineNumber	58189.0	1.887453	1.018829	1.0000
OrderQuantity	58189.0	1.569386	1.047532	1.0000
UnitPrice	58189.0	413.888218	833.052938	0.5725
TotalProductCost	58189.0	296.539185	560.171436	0.8565
SalesAmount	58189.0	503.666270	941.462817	2.2900
TaxAmt	58189.0	40.293303	75.317027	0.1832
Unnamed: 13	0.0	NaN	NaN	NaN
Unnamed: 14	0.0	NaN	NaN	NaN
Unnamed: 15	58189.0	503.666269	941.462815	2.2900
Unnamed: 16	58189.0	0.000001	0.000014	0.0000
Unnamed: 17	0.0	NaN	NaN	NaN
Unnamed: 18	58189.0	38.398254	667.349417	-5106.9068
Unnamed: 19	0.0	NaN	NaN	NaN
StandardCost_x	58189.0	296.539185	560.171436	0.8565
List Price	58189.0	503.666270	941.462817	2.2900
Unnamed: 22	0.0	NaN	NaN	NaN
diif std cost	58189.0	0.000000	0.000000	0.0000
diff list price	58189.0	0.000000	0.000000	0.0000

StandardCost_y	58189.0	296.539185	560.171436	0.8565
ListPrice	58189.0	503.666270	941.462817	2.2900
DaysToManufacture	58189.0	1.045215	1.757395	0.0000
YearlyIncome	58189.0	59769.887779	33128.041818	10000.0000
TotalChildren	58189.0	1.838921	1.614467	0.0000
NumberChildrenAtHome	58189.0	1.073502	1.580055	0.0000
HouseOwnerFlag	58189.0	0.690560	0.462267	0.0000
NumberCarsOwned	58189.0	1.502466	1.155496	0.0000
	25%	50%	75%	max
SalesOrderLineNumber	1.0000	2.0000	2.0000	8.0000
OrderQuantity	1.0000	1.0000	2.0000	4.0000
UnitPrice	4.9900	24.4900	269.9950	3578.2700
TotalProductCost	3.3623	12.1924	343.6496	2171.2942
SalesAmount	8.9900	32.6000	539.9900	3578.2700
TaxAmt	0.7192	2.6080	43.1992	286.2616
Unnamed: 13	NaN	NaN	NaN	NaN
Unnamed: 14	NaN	NaN	NaN	NaN
Unnamed: 15	8.9900	32.6000	539.9900	3578.2700
Unnamed: 16	0.0000	0.0000	0.0000	0.0003
Unnamed: 17	NaN	NaN	NaN	NaN
Unnamed: 18	1.4335	6.2537	21.9037	1487.8356
Unnamed: 19	NaN	NaN	NaN	NaN
StandardCost_x	3.3623	12.1924	343.6496	2171.2942
List Price	8.9900	32.6000	539.9900	3578.2700
Unnamed: 22	NaN	NaN	NaN	NaN
diif std cost	0.0000	0.0000	0.0000	0.0000

diff list price	0.0000	0.0000	0.0000	0.0000
StandardCost_y	3.3623	12.1924	343.6496	2171.2942
ListPrice	8.9900	32.6000	539.9900	3578.2700
DaysToManufacture	0.0000	0.0000	4.0000	4.0000
YearlyIncome	30000.0000	60000.0000	80000.0000	170000.0000
TotalChildren	0.0000	2.0000	3.0000	5.0000
NumberChildrenAtHome	0.0000	0.0000	2.0000	5.0000
HouseOwnerFlag	0.0000	1.0000	1.0000	1.0000
NumberCarsOwned	1.0000	2.0000	2.0000	4.0000

Standard deviation is square root of variance.

```
df.duplicated().sum()
0

def missing_pct(df):
    # Calculate missing value and their percentage for each column
    missing_count_percent = df.isnull().sum() * 100 / df.shape[0]
    df_missing_count_percent =
pd.DataFrame(missing_count_percent).round(2)
    df_missing_count_percent =
df_missing_count_percent.reset_index().rename(
    columns={
        'index': 'Column',
        0: 'Missing_Percentage (%)'
    }
)
df_missing_value = df.isnull().sum()
df_missing_value = df_missing_value.reset_index().rename(
    columns={
        'index': 'Column',
        0: 'Missing_value_count'
    }
)
# Sort the data frame
#df_missing = df_missing.sort_values('Missing_Percentage (%)',
ascending=False)
Final = df_missing_value.merge(df_missing_count_percent, how =
'inner', left_on = 'Column', right_on = 'Column')
Final = Final.sort_values(by = 'Missing_Percentage (%)', ascending
```

```
= False)
    return Final
```

```
missing_pct(df)
```

	Column	Missing_value_count	Missing_Percentage (%)
22	Unnamed: 22	58189	100.00
19	Unnamed: 19	58189	100.00
14	Unnamed: 14	58189	100.00
13	Unnamed: 13	58189	100.00
17	Unnamed: 17	58189	100.00
30	Color	27442	47.16
0	ProductKey	0	0.00
42	MaritalStatus	0	0.00
41	BirthDate	0	0.00
39	LastName	0	0.00
40	FullName	0	0.00
38	FirstName	0	0.00
37	StartDate	0	0.00
36	ProductDescription	0	0.00
35	Photo	0	0.00
34	ModelName	0	0.00
43	Gender	0	0.00
44	YearlyIncome	0	0.00
32	DaysToManufacture	0	0.00
45	TotalChildren	0	0.00
46	NumberChildrenAtHome	0	0.00
47	Education	0	0.00
48	Occupation	0	0.00
49	HouseOwnerFlag	0	0.00
50	NumberCarsOwned	0	0.00
51	AddressLine1	0	0.00
52	DateFirstPurchase	0	0.00
53	CommuteDistance	0	0.00
54	Region	0	0.00
55	Country	0	0.00
56	Group	0	0.00
33	ProductLine	0	0.00
29	StandardCost_y	0	0.00
31	ListPrice	0	0.00
12	TaxAmt	0	0.00
2	ShipDate	0	0.00
3	CustomerKey	0	0.00
4	PromotionKey	0	0.00
5	SalesTerritoryKey	0	0.00
6	SalesOrderNumber	0	0.00
7	SalesOrderLineNumber	0	0.00
8	OrderQuantity	0	0.00
9	UnitPrice	0	0.00
10	TotalProductCost	0	0.00

11	SalesAmount	0	0.00
15	Unnamed: 15	0	0.00
1	OrderDate	0	0.00
16	Unnamed: 16	0	0.00
18	Unnamed: 18	0	0.00
20	StandardCost_x	0	0.00
21	List Price	0	0.00
23	diif std cost	0	0.00
24	diff list price	0	0.00
25	DateKey	0	0.00
26	ProductName	0	0.00
27	SubCategory	0	0.00
28	Category	0	0.00
57	RegionImage	0	0.00

```
df= df.dropna(axis=1)
```

```
# Extracting Year from OrderDate
```

```
df['sale_year'] = df['OrderDate'].dt.year
```

```
# Extracting Month from OrderDate
```

```
df['sale_month'] = df['OrderDate'].dt.month
```

```
# Extracting day from OrderDate
```

```
df['sale_day'] = df['OrderDate'].dt.day
```

```
# Extracting dayofweek from OrderDate
```

```
df['sale_week'] = df['OrderDate'].dt.dayofweek
```

```
# Extracting day_name from OrderDate
```

```
df['sale_day_name'] = df['OrderDate'].dt.day_name()
```

```
# Extracting Month Year from OrderDate
```

```
df['year_month'] = df['OrderDate'].apply(lambda x:x.strftime('%Y-%m'))
```

```
# Calculate Total Invoice Amount
```

```
df['total_invoice_amount'] = df['SalesAmount'] + df['TaxAmt']
```

```
# Considering only salesamount and total_sales_amount to calculate profit
```

```
df['profit'] = (df['UnitPrice']*df['OrderQuantity']) -  
df['TotalProductCost']
```

```
# Removing extra character from the string
```

```
df['ProductName'] = df['ProductName'].str.replace(',', '-')
```

```
# Calculate Age
```


```
df['Age'] = df['OrderDate'].dt.year - df['BirthDate'].dt.year
```

```
df['Category'].unique().tolist()
```

```
df['SubCategory'].unique().tolist()
```

```
['Road Bikes',  
'Mountain Bikes',  
'Bottles and Cages',  
'Gloves',  
'Tires and Tubes',  
'Helmets',  
'Touring Bikes',  
'Jerseys',  
'Cleaners',  
'Caps',  
'Hydration Packs',  
'Socks',  
'Fenders',  
'Vests',  
'Bike Racks',  
'Bike Stands',  
'Shorts']
```

```
Avg_unit_price = df.groupby(['ProductKey'])['UnitPrice'].mean()  
ax = sns.distplot(Avg_unit_price, kde=True, hist=True)  
ax.set(title='Distribution of Average unit price',  
       xlabel='Average Unit Price');
```




Unit product price is maximum between \$0 \$ to1000.




```
n_orders = df.groupby(['CustomerKey'])['SalesOrderNumber'].nunique()
multi_orders_perc = np.sum(n_orders > 1)/df['CustomerKey'].nunique()
print(f"{100*multi_orders_perc:.2f}% of customers ordered more than once.")
```

36.97% of customers ordered more than once.

```
ax = sns.distplot(n_orders)
ax.set(title='Distribution of number of orders per customer',
       xlabel='# of orders',
       ylabel='# of customers');
```




```
n_salesordernumber = df.groupby(['SalesOrderNumber'])
['SalesOrderLineNumber'].transform('max')
ax = sns.distplot(n_salesordernumber, kde=False, color='#374045')
ax.set(title='Distribution of sales order line number',
       xlabel='# of Sales order line number',
       ylabel='# of orders');
```



Three to two products are ordered in a single order most of the time.

```
n_order_quantity = df.groupby(['SalesOrderNumber'])  
['OrderQuantity'].sum()  
ax = sns.distplot(n_order_quantity, kde=True, hist=True)  
ax.set(title='Distribution of order_quantity',  
        xlabel='# of order_quantity',  
        );
```




Maximum quantity ordered for a product is below 5

```
bins = [18, 30, 40, 50, 60, 70, 120]
labels = ['18-29', '30-39', '40-49', '50-59', '60-69', '70+']
df['agerange'] = pd.cut(df.Age, bins, labels = labels, include_lowest =
True)

age_distribution =
df['agerange'].value_counts().to_frame().reset_index()


age_distribution.columns = ['Age Range', 'Population count']

sns.lineplot( x='Age Range', y='Population
count', data=age_distribution)
plt.show()
```



A sizable portion of the clientele is made up of people between the ages of 40 and 59.

```
df.groupby('sale_year')['SalesAmount'].sum().plot(kind='bar',  
color='#374045');
```




The year 2016 saw an exponential surge in sales.

```
top_selling_product = df.groupby(['Category', 'SubCategory',  
'ProductName'])['OrderQuantity'].sum().nlargest(5).to_frame()  
top_selling_product
```

			OrderQuantity
Category	SubCategory	ProductName	
Accessories	Bottles and Cages	Water Bottle - 30 oz.	6370
		Patch Kit/8 Patches	4705
		Mountain Tire Tube	4551
	Tires and Tubes	Road Tire Tube	3544
		Sport-100 Helmet- Red	3398

```
top_selling_product.reset_index(inplace=True)  
sns.barplot(x='ProductName',  
y='OrderQuantity',data=top_selling_product)  
plt.xticks(rotation=90)  
plt.show()
```




```

cat_subcat_qty = df.groupby(['sale_year', 'Category', 'SubCategory'])
['OrderQuantity'].sum().to_frame()
cat_subcat_qty = cat_subcat_qty.sort_values(['sale_year', 'Category'],
ascending=True)
cat_subcat_qty.style.bar(subset=['OrderQuantity'])

<pandas.io.formats.style.Styler at 0x2024bb47e50>

country_qty_sales = df.groupby('Country')
['OrderQuantity'].sum().sort_values(ascending=False)
country_qty_sales.plot(kind='bar');

```



High quantity of products is ordered from Australia and United States

```

cat_subcat_profit = df.groupby(['sale_year', 'Category',
'SubCategory'])['profit'].sum().to_frame()

#Sorting the results
cat_subcat_profit = cat_subcat_profit.sort_values(['sale_year',
'Category'], ascending=True)
cat_subcat_profit.style.bar(subset=['profit'])

<pandas.io.formats.style.Styler at 0x202470c9a00>

```


Major Profit is contributed by the Bike Category

```
df.groupby(['Category', 'SubCategory', 'ProductName'])  
['profit'].sum().nsmallest(10).to_frame()
```

			profit
Category	SubCategory	ProductName	
Clothing	Socks	Racing Socks- L	1474.4574
		Racing Socks- M	1581.3837
Accessories	Cleaners	Bike Wash - Dissolver	4299.8688
	Tires and Tubes	Patch Kit/8 Patches	4314.8350
Clothing	Caps	AWC Logo Cap	4331.8315
Accessories	Tires and Tubes	Touring Tire Tube	4363.8089
Clothing	Jerseys	Long-Sleeve Logo Jersey- XL	4495.6007
		Short-Sleeve Classic Jersey- L	4544.8782
		Long-Sleeve Logo Jersey- S	4610.5777
		Short-Sleeve Classic Jersey- M	4793.2322


```
country_sales = pd.DataFrame(df.groupby('Country').sum()  
[['SalesAmount', 'profit']])  
country_sales.reset_index(inplace=True)
```

```
sns.barplot(data=country_sales, x='Country', y='profit')  
plt.xticks(rotation=30)  
plt.show()
```



High volume of profit is earned from Australia and United States


```
df['OrderreadyDate'] = df['OrderDate'] +  
pd.to_timedelta(df['DaysToManufacture'], unit='D')  
  
# Check the delay between order shipment date and order ready to  
supply  
df['shipping_efficiency'] = (df['ShipDate'] -  
df['OrderreadyDate']).dt.days  
  
plt.hist(data=df, x="shipping_efficiency",)  
plt.show()
```



The average order has a gap of 7 days between the day the order is ready for export from the factory and the date it was shipped. Management must work to reduce this gap toward 3 days.

```
dj = df.groupby('agerange')['SalesAmount'].mean().to_frame().dropna()  
dj.reset_index(inplace=True)  
sns.barplot(data=dj, x='agerange', y='SalesAmount')  
plt.show()
```





40-49 age group has produced most revenue.


```
def create_bins(lower_bound, width, quantity):
    """ create_bins returns an equal-width (distance) partitioning.
        It returns an ascending list of tuples, representing the
        intervals.
        A tuple bins[i], i.e. (bins[i][0], bins[i][1]) with i > 0
        and i < quantity, satisfies the following conditions:
        (1) bins[i][0] + width == bins[i][1]
        (2) bins[i-1][0] + width == bins[i][0] and
            bins[i-1][1] + width == bins[i][1]
    """

    bins = []
    for low in range(lower_bound,
                     lower_bound + quantity*width + 1, width):
        bins.append((low, low+width))
    return bins

bins = create_bins(lower_bound=10000,
                   width=30000,
                   quantity=5)
bins2 = pd.IntervalIndex.from_tuples(bins)
df['salary_range'] = pd.cut(df['YearlyIncome'], bins2)


df_4 = df.groupby('salary_range')['SalesAmount'].mean().to_frame()
df_4.reset_index(inplace=True)
```

```
sns.barplot(x="SalesAmount", y="salary_range", data=df_4)
plt.show()
```

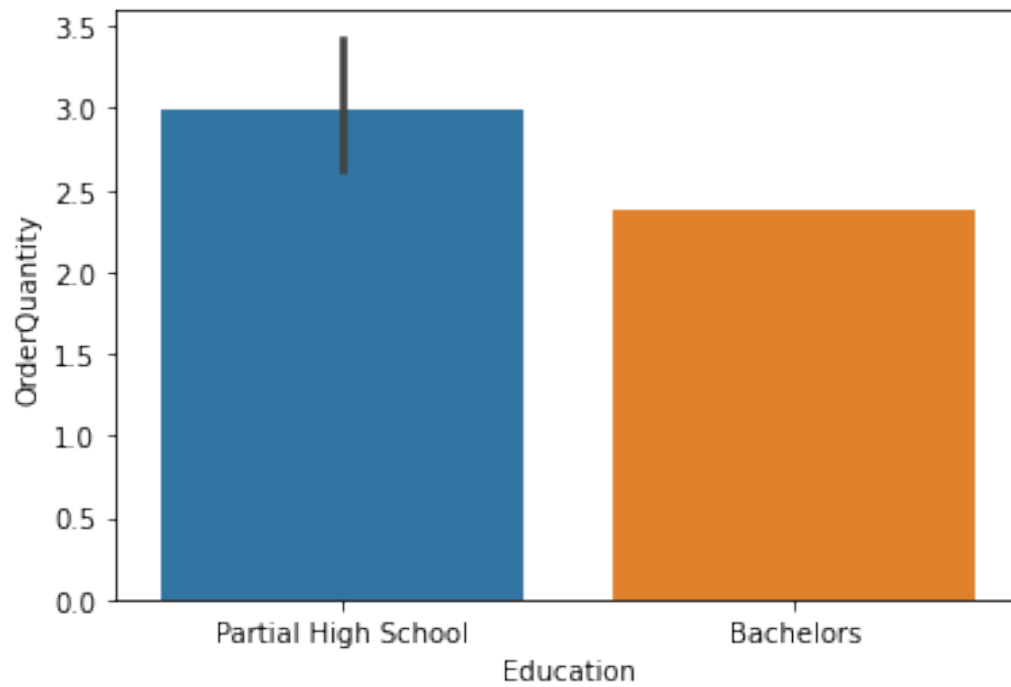


High salary range leads to increase in purchase

```
df_6 = df[(df['Education']=='Partial High School')|
(df['Education']=='Bachelors')].groupby('Education')
['YearlyIncome'].mean().to_frame()
df_6.reset_index(inplace=True)
sns.barplot(data=df_6, x='Education', y='YearlyIncome')
plt.show()
```



```
df_7 = df[(df['Education']=='Partial High School')|  
(df['Education']=='Bachelors')]  
df_7 = df_7.groupby(['Education', 'ProductName'])  
['OrderQuantity'].mean().to_frame().sort_values('OrderQuantity',  
ascending=False)[:10]  
df_7.reset_index(inplace=True)  
sns.barplot(data=df_7, x="Education",  
            y="OrderQuantity")  
plt.show()
```



Customers with a high school diploma and modest annual income buy more products than people with bachelor's degrees.